

# Data Level Parallelism: SIMD Vector Architecture, Graphics Processing Units, Systolic Arrays. Interconnection Networks

**Prasanta K. Jana, IEEE Senior Member**





**Department of Computer Science and Engineering  
Indian Institute of Technology (ISM), Dhanbad  
E-mail: [prasantajana@yahoo.com](mailto:prasantajana@yahoo.com)**

# SIMD Computer Organisation



Each  $PE_i$  is arithmetic logic unit (ALU).  
Each  $PME_i$  is local storage.





- Formally, an SIMD computer  $C$  is characterized by the following set of parameters:


$$C = \langle N, F, I, M \rangle$$

where  $N$  = the no. of  $PE$ s in the system.

$F$  = a set of data-routing functions.

$I$  = the set of machine instructions.

$M$  = the set of masking schemes.



Components in a Processing Element( $PE_i$ )

# An Illustration of Data routing and masking schemes



- Let  $A = (A_0, A_1, \dots, A_{n-1})$ , we need to compute the following  $n$  summations:


$$s(k) = \sum_{i=0}^k A_i, \text{ for } k = 0, 1, \dots, n-1$$

- These  $n$  vector summations can be computed recursively by :

$$s(0) = A_0,$$

$$s(k) = s(k-1) + A_k, \text{ for } k = 0, 1, \dots, n-1$$

# SIMD Computer Organisation



The calculation of the summation  $s(k) = \sum_{i=0}^k A_i$ , for  $k = 0, 1, \dots, 7$  in an SIMD Machine


- Routing function:  $R_i(j) = j + 2^i, 0 \leq i \leq \log(n) - 1, 0 \leq j \leq n - 1$



- Static versus Dynamic networks.

The SIMD interconnection networks are classified into the following two categories based on network topologies: Static networks versus dynamic networks.


# SIMD Interconnection Networks




Static interconnection network topologies



- We consider two classes of dynamic networks:
  - Single-stage
  - Multi-stage



Conceptual view of a single-stage interconnection network



A two-by-two switching box and its four interconnection states.

# Evaluation criteria of Interconnection network



- **Diameter:** Maximum of shortest distance between any two arbitrary nodes.
- **Bisection width:** The minimum number of edges to be removed to divide the network into two halves.
- **Number of edges per node.**
- **Maximum edge length.**

- Dynamic networks can be divided into three classes:
  - Blocking
  - Rearrangeable
  - Nonblocking

- **Blocking networks:** Simultaneous connections of more than one terminal pair may result in conflicts. Examples of a blocking network are data manipulator, omega, flip, n cube, and baseline.
- **Rearrangeable networks:** It can perform all possible connections between inputs and outputs by rearranging its existing connections so that a connection path for a new input-output pair can always be established. Example is Benes network.





- $$C_i(a_{n-1} \dots a_1 a_0) = a_{n-1} \dots a_{i+1} \overline{a_i} a_{i-1} \dots a_0, \text{ for } i = 0, 1, \dots, n-1$$

# Implementation as a single stage network



Routing Function



(b) The recirculating cube network

- Shuffle function is defined as follows:

$$S(a_{n-1} \dots a_1 a_0) = a_{n-2} \dots a_1 a_0 a_{n-1}$$

- Inverse shuffle function is defined as follows:

$$IS(a_{n-1} \dots a_1 a_0) = a_0 a_{n-1} \dots a_1$$


- Exchange routing function is defined as follows:

$$E(a_{n-1} \dots a_1 a_0) = a_{n-1} \dots a_1 \overline{a_0}$$


- Note that  $E(A) = C_0(A)$

- Butterfly function:  $\beta_k(b_{n-1} b_{n-2} \dots b_{k+1} b_k b_{k-1} \dots b_1 b_0) = b_{n-1} b_{n-2} \dots b_{k+1} b_0 b_{k-1} \dots b_1 b_k$






Single stage shuffle exchange network (solid lines are exchange and dashed lines are shuffle).




Shuffle connection

# Multi-stage implementation




A Multistage cube network for  $N = 8$

# Multi-stage implementation



Three stage  $8 \times 8$  Omega switching network



The Omega network with switch box repositioned

- The cube network uses two function switch boxes, whereas the omega network uses four function switch boxes.
- The data flow directions in the two networks are opposite to each other. In other words, the roles of input-output lines are exchanged in the two networks.




$$R_{+1}(i) = (i + 1) \bmod N$$

$$R_{-1}(i) = (i - 1) \bmod N$$


$$R_{+r}(i) = (i + r) \bmod N$$

$$R_{-r}(i) = (i - r) \bmod N$$



An Illiac network with  $N = 16$  and  $r = 4$





An Illiac network with  $N = 16$  PEs

- The permutation cycles  $(a, b, c), (d, e)$  stands for the permutation  $a \rightarrow b, b \rightarrow c, c \rightarrow a$ , and  $d \rightarrow e, e \rightarrow d$  in a circular fashion within a pair of parenthesis:

$$R_{+1} = (0 \ 1 \ 2 \ \dots \ N - 1)$$

$$R_{-1} = (N - 1 \ \dots \ 2 \ 1 \ 0)$$

$$R_{+r} = \prod_{i=0}^{r-1} (i \ i + r \ i + 2r \ \dots \ i + N - r)$$

$$R_{-r} = \prod_{i=0}^{r-1} (i + N - r \ \dots \ i + 2r \ i + r \ i)$$

- Routing functions:

$$B_{+j}(i) = (i + 2^j)(\text{mod } N)$$

$$B_{-j}(i) = (i - 2^j)(\text{mod } N)$$

where  $0 \leq i \leq N - 1$ ,  $0 \leq j \leq N - 1$  and  $n = \log_2 N$

$$B_{+0} = R_{+1}$$

$$B_{-0} = R_{-1}$$

$$B_{+n/2} = R_{+r}$$

$$B_{-n/2} = R_{-r}$$

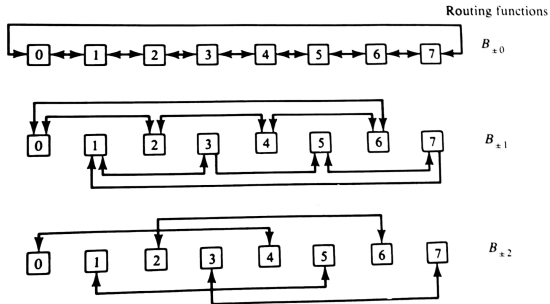
$$B \leq (\log_2 N)/2$$

$$B_{\pm 0}(j) = (j \pm 1) \bmod 8$$

$$B_{\pm 1}(j) = (j \pm 2) \bmod 8$$

$$B_{\pm 2}(j) = (j \pm 4) \bmod 8$$

where  $N = 8$  and  $0 \leq j \leq 7$



A recirculating barrel shifter for  $N = 8$ .


$$S_N = \frac{\sqrt{N}-1}{(\log_2 N)/2} = \frac{2^k-1}{k}$$

# Systolic System




- Consist of a set of inter-connected cells
- Each capable of performing some simple operation

- By replacing a simple PE with an array of PEs, a higher computation throughput can be achieved (as shown in figure) without increasing memory.



Conventional Processor




A Systolic Processor Array







- The function is analogous to that of a heart, that is
  - It 'pulses' data through the array of PEs. Once data item is brought out from the memory.
  - It can be used effectively at each cell, it passes.
  - This is possible for a wide class of computation-bound computation.




Fir-filter, Convolution,DFT etc.




Mesh for DP, Graph Algorithm Matrix-multi



Matrix multi; (banded), LU decomposition  
**Two dimensional hexagonal array**



Tree Searching , Recurrence evaluation etc.



**Triangular array**  
Inversion of triangular matrix

# Graeffe's Root Squaring



$$C_j = A_j^2 - 2A_{j-1}A_{j+1} + 2A_{j-2}A_{j+2} - 2A_{j-3}A_{j+3} + 2A_{j-4}A_{j+4} - \dots$$

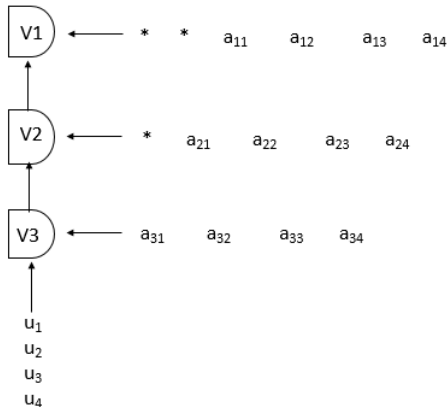


$A_0$	$A_1$	$A_2$	$A_3$	$A_4$	$A_5$	$\rightarrow$
$A_0$	$A_1$	$A_2$	$A_3$	$A_4$	$A_5$	$\leftarrow$
0	$A_0$	$A_1$	$A_2$	$A_3$	$A_4$	
$A_1$	$A_2$	$A_3$	$A_4$	$A_5$	0	

# Matrix Vector Multiplication



$$[A]_{3 \times 4} \times [U]_{4 \times 1} = [V]_{3 \times 1}$$



$$\begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{11} & a_{22} & a_{23} & a_{24} \\ a_{11} & a_{11} & a_{11} & a_{11} \end{pmatrix} \begin{pmatrix} a_{11} \\ a_{11} \\ a_{11} \\ a_{11} \end{pmatrix}$$

- Convolution : Given  $\{w_1, w_2, \dots, w_n\}$  and an input sequence  $\{x_1, x_2, \dots, x_n\}$

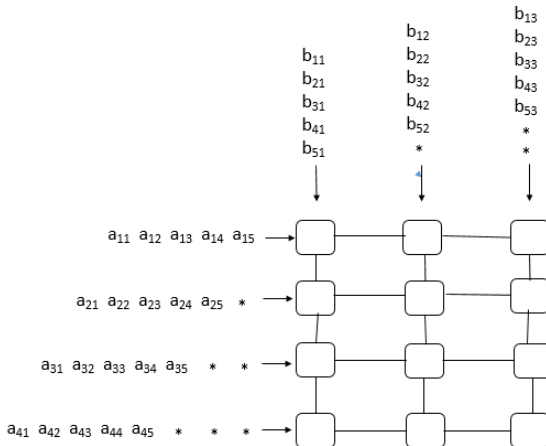
Compute  $\{y_1, y_2, \dots, y_{2n-1}\}$

by  $y_i = \sum_{j=1}^n x_{i-j+1} \times w_j$  where  $1 \leq i \leq 2n - 1$

For  $n = 3$ , this is equivalent or

$$\begin{bmatrix} x_1 & 0 & 0 \\ x_2 & x_1 & 0 \\ x_3 & x_2 & x_1 \\ 0 & x_3 & x_2 \\ 0 & 0 & x_3 \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ y_5 \end{bmatrix}$$

# Matrix-Matrix Multiplication





# Matrix-Matrix Multiplication



Procedure MESH-MUL(A,B,C)  $(A)_{m \times n}, (B)_{n \times k}$   
for i=1 to m do in parallel  
  for j=1 to k do in parallel  
    (1)  $e_{ij} \leftarrow 0$   
    (2) while P (i, j) receives two inputs a & b do  
      (i)  $e_{ij} \leftarrow e_{ij} + a \times b$   
      (ii) if  $i < m$  then send b to P(i + 1, j) endif  
      (iii) if  $j < k$  then send a to P (i, j + 1) endif  
    endwhile  
  endfor  
endfor


- Time Complexity : Elements  $a_{m1}$  and  $b_{1k}$  take  $m + k + n - 2$  steps from the beginning to reach P(m, k), the last processor to terminal.  
Assuming  $n \geq m$  and  $n \geq k$   
it requires  $O(n)$  time.

# Palindrome Recognition



- Palindrome is a string of the form  $ww^r$  where  $w^r$  is the reverse of  $w$ . For example, 110011 is a palindrome, but 110001, 101 are not.  
Input :  $x_1, x_2, \dots, x_w$   
Output : Whether or not  $x_1, x_2, \dots, x_i$  is a palindrome immediately after inputting  $x_i$ .

# Palindrome Recognition



- **Checking** : The string  $x_1, x_2, \dots, x_i$  is a palindrome ,iff, each cell contains identical values at the  $i^{th}$  step.



- Definition: A Graphics Processing Unit (GPU) is a single-chip processor primarily used to manage and boost the performance of video and graphics. GPU features include:
- 2-D or 3-D graphics
- Digital output to flat panel display monitors
- Application support for high-intensity graphics software such as AutoCAD
- Used in a PC on a video card or motherboard; it is also used in mobile phones, display adapters, workstations and game consoles.
- These features are designed to lessen the work of the CPU and produce faster video and graphics.



- The first GPU was developed by NVidia in 1999 and called the GeForce 256.
- This GPU model could process 10 million polygons per second and had more than 22 million transistors.
- Generally the GPU is connected to the CPU and is completely separate from the motherboard.
- The random access memory (RAM) is connected through the accelerated graphics port (AGP) or the peripheral component interconnect express (PCI-Express) bus.
- Most GPUs use their transistors for 3-D computer graphics.
- Applications such as computer-aided design (CAD) can process over 200 billion operations per second



- GPU-accelerated computing is the use of a GPU together with a CPU to accelerate deep learning, analytics, and engineering applications.
- GPU-accelerated computing offloads compute-intensive portions of the application to the GPU, while the remainder of the code still runs on the CPU. From a user's perspective, applications simply run much faster.



- A simple way to understand the difference between a GPU and a CPU is to compare how they process tasks.
- A CPU consists of a few cores optimized for sequential serial processing while a GPU has a massively parallel architecture consisting of thousands of smaller, more efficient cores designed for handling multiple tasks simultaneously.
- GPUs have thousands of cores to process parallel workloads efficiently





THANK YOU